GPU Programming, Spring 2019

Patric Manner (37208)

February 6, 2019

Hardware

Asus N56VZ laptop with a Nvidia GeForce GT 650M graphics card; 384 CUDA cores and a compute capability of 3.0 (Kepler architecture).

Design and execution

General

- Removed the 100000 from the beginning of each file

CPU side

- Allocated unified memory with cudaMallocManaged()
- Read values from the files and converted to radians into 4 arrays (ascension and declination separately for both files)
- Initialize the histogram arrays with zeros

GPU side

- Threads in block: 512
- Blocks in grid: 196
- Check if index ¡ N
- Each thread multiplies one value from the first galaxy list (it's index) with all of the values in the second galaxy list
- Use given formula
- Convert to degrees, determine bin number, add to bin with atomicAdd

Program specs and results

- Threads in block: 512
- Blocks in grid: 196
- Average running time: 36s
- Registers used: 22
- Omega results point towards the real galaxies being non-randomly distributed

Problems encountered

- Got bus errors due to the acos function getting input values greater than 1, which led to that the program tried to access a negative array index in the histogram
- Took a long time to realize that I was printing %d even though the result of the aggregation was correct when checking if the histograms were filled properly